

# Clustering

DSE 220

# So far..

- Introduction
- Nonparametric Methods
  - Decision Trees
  - kNN
- Parametric Methods
  - Generative Models
    - Naive Bayes
    - Binary Features, Multinomial Features
    - Gaussian Generative Model
    - Fisher Linear Discriminant Analysis
- Regression Analysis
- Evaluation Metrics
  - Accuracy, Recall, Precision, Sensitivity, Specificity,...
  - ROC Curves

# So far..

- Parametric Methods
  - Discriminative Methods
    - Logistic Regression
      - Gradient Descent
      - Newton-Raphson
    - Perceptron
    - SVMs
  - Kernels
  - Richer Output Spaces

# Today

- Representation Learning
  - k-means
  - EM
  - Agglomerative hierarchical clustering
  - ***Hands-On / Self – practice***
- Informative Projections
  - PCA
  - SVD
  - ***Hands-On / Self – practice***
- Beyond Projections
  - Manifold Learning
  - Dictionary Learning
  - ***Hands-On / Self – practice***
- Ensemble Methods
  - Boosting
  - Bagging
  - Random Forests
  - ***Hands-On / Self – practice***

# Representation learning



Good representations make learning easier.

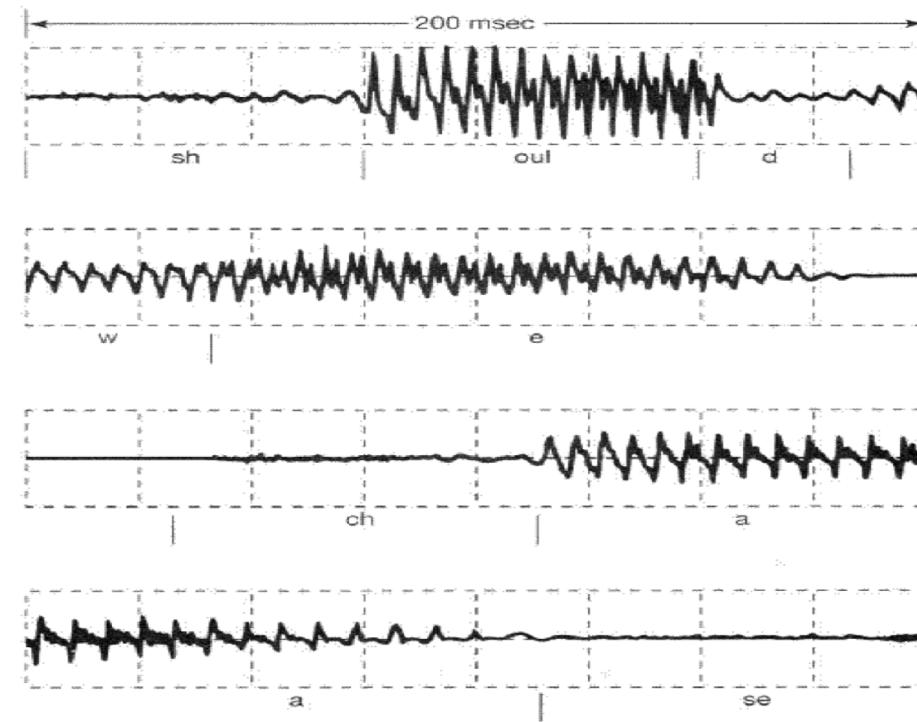
# Representation learning



Good representations make learning easier.

- They bring out the true degrees of freedom in the data.
- They capture relevant structure at multiple scales.
- They screen out noisy or irrelevant structure.

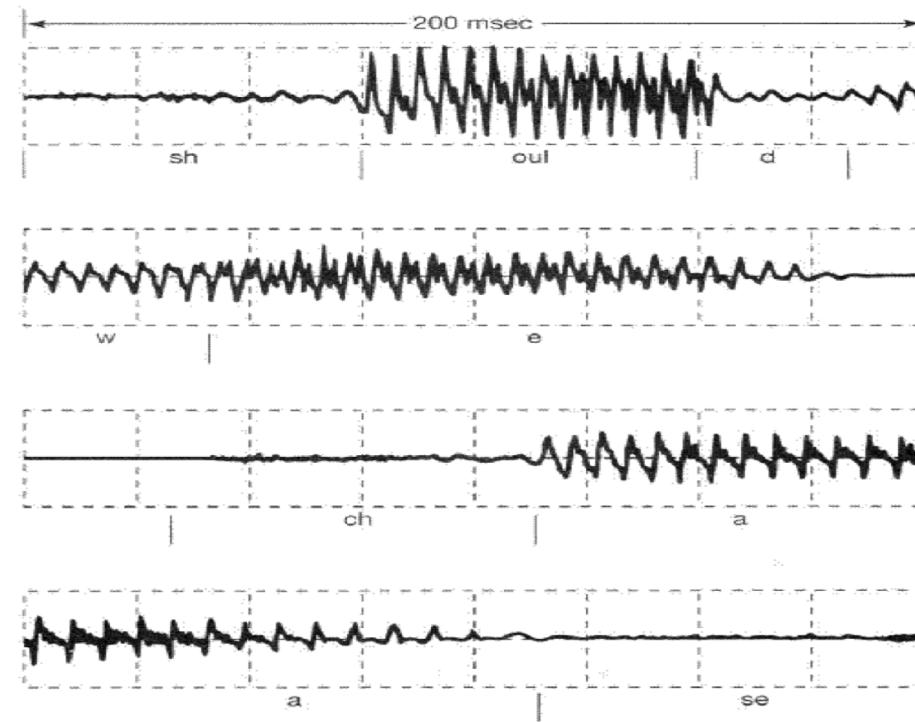
# Degrees of freedom



Usual representation of speech:

- Take overlapping windows of the speech signal
- Apply many filters within each window
- More filters  $\Rightarrow$  higher dimensional

# Degrees of freedom



Usual representation of speech:

- Take overlapping windows of the speech signal
- Apply many filters within each window
- More filters  $\Rightarrow$  higher dimensional

But the speech is produced by a physical system (vocal tract) with a fixed number of degrees of freedom. And the phoneme being uttered can be characterized by the configuration of this apparatus.

# Multiscale structure



Commonly-occurring structure at many levels:

- Low-level: like local edges
- Higher-level: like wheels, windows

# Representation learning: goals

How, and to what extent, can underlying degrees of freedom and multiscale structure be learned from the statistics of unlabeled data?

# Representation learning: goals

How, and to what extent, can underlying degrees of freedom and multiscale structure be learned from the statistics of unlabeled data?

And when labels are available, how can a good representation be learned in tandem with the classifier?

# Representation learning: goals

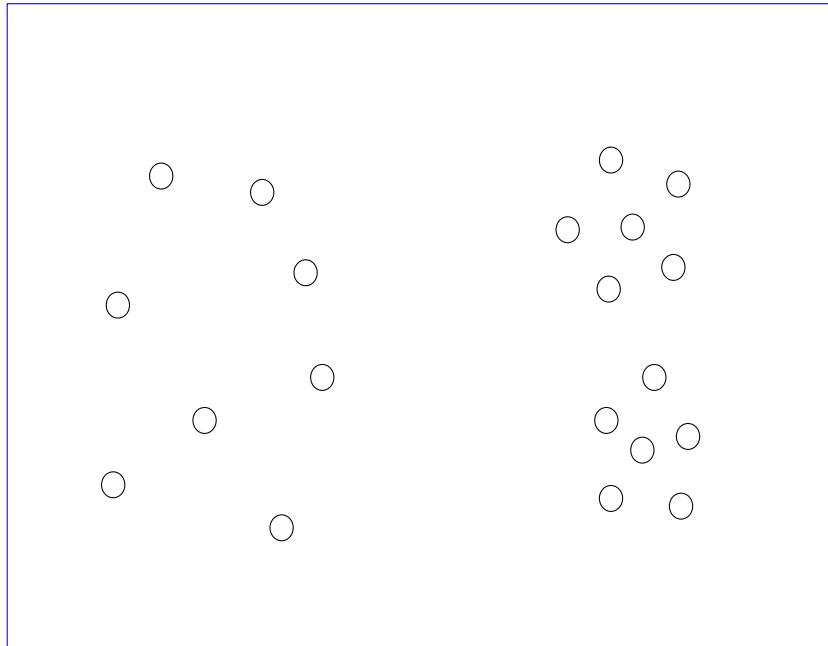
How, and to what extent, can underlying degrees of freedom and multiscale structure be learned from the statistics of unlabeled data?

And when labels are available, how can a good representation be learned in tandem with the classifier?

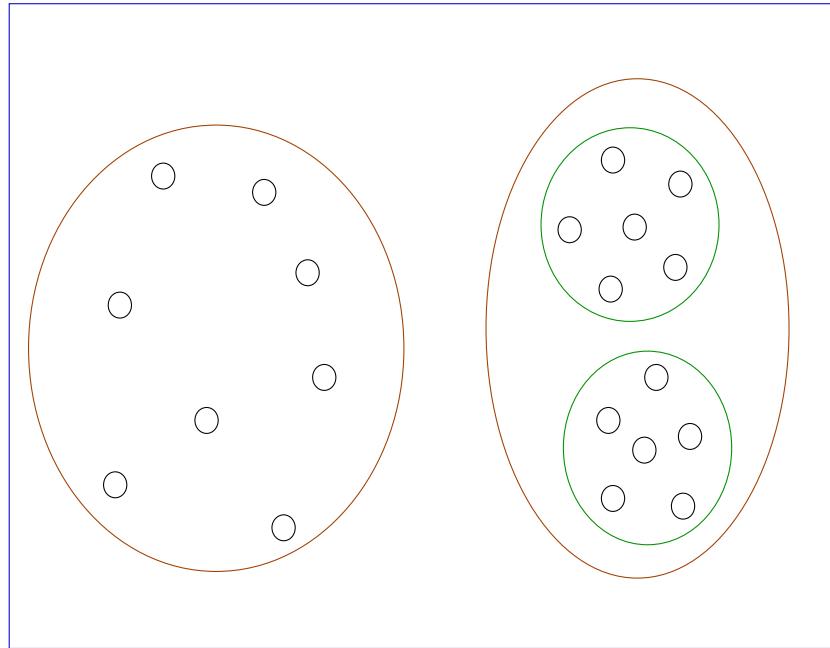
Topics:

- Clustering
- Informative linear projections
- Embedding and manifold learning
- Metric learning
- Autoencoders
- Deep nets

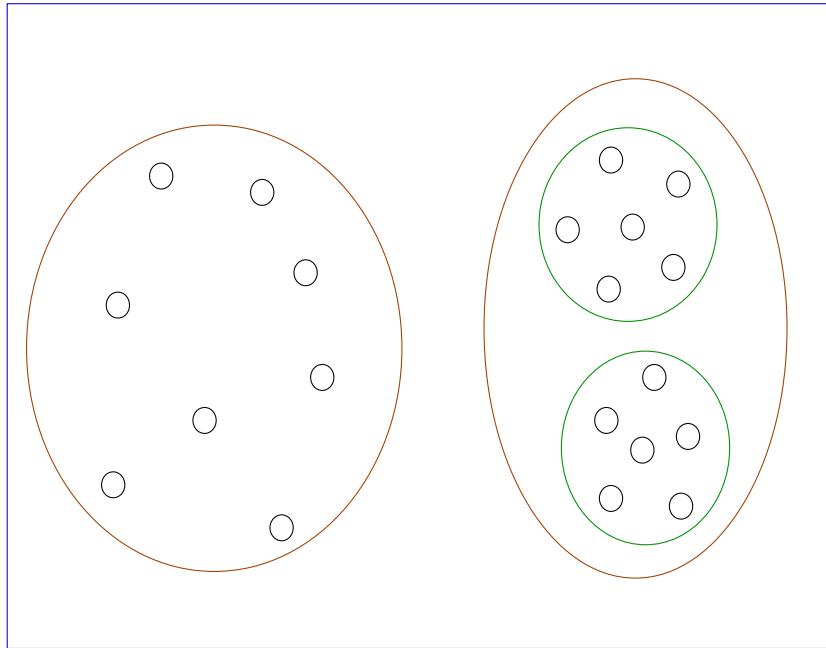
# Clustering in $\mathbb{R}^p$



# Clustering in $\mathbb{R}^p$

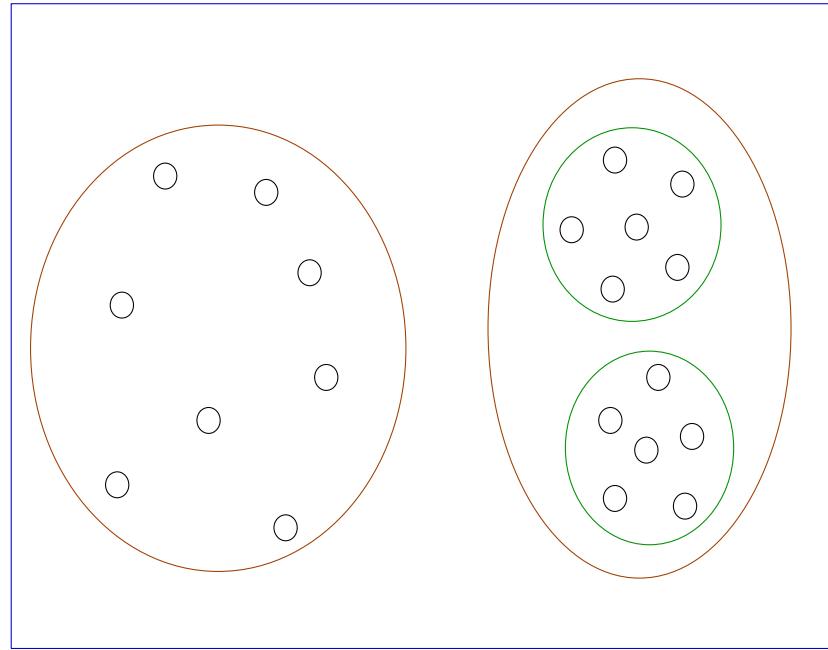


# Clustering in $\mathbb{R}^p$



Two common uses of clustering:

# Clustering in $\mathbb{R}^p$



Two common uses of clustering:

- **Vector quantization**  
Find a finite set of representatives, that provides good coverage of a complex, possibly infinite, high-dimensional space
- **Finding meaningful structure in data**  
Finding salient grouping in data.

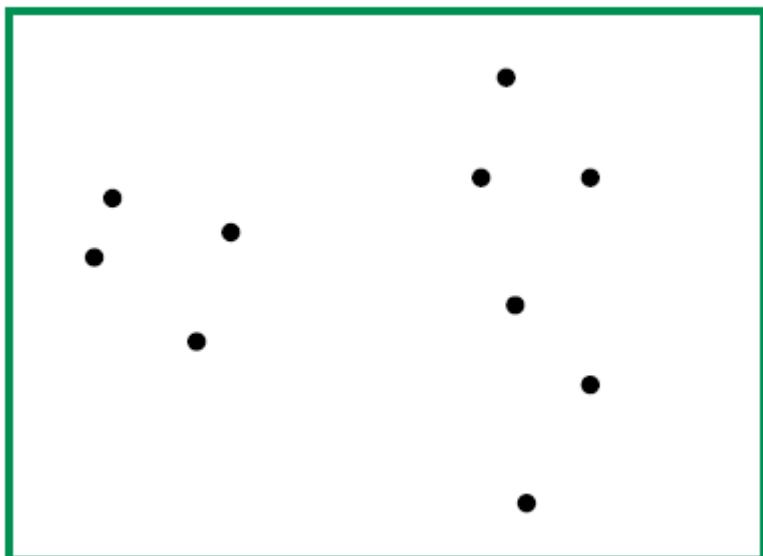
# Widely-used clustering methods

- ①  $K$ -means and its many variants
- ② EM for mixtures of Gaussians
- ③ Agglomerative hierarchical clustering

# The $k$ -means optimization problem

- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^p$ ; integer  $k$
- Output: “Centers”, or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:

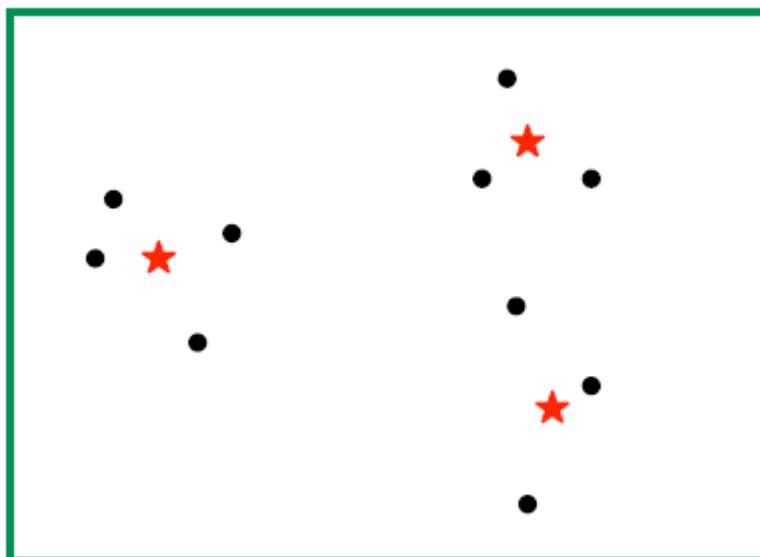
$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



# The $k$ -means optimization problem

- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^p$ ; integer  $k$
- Output: “Centers”, or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:

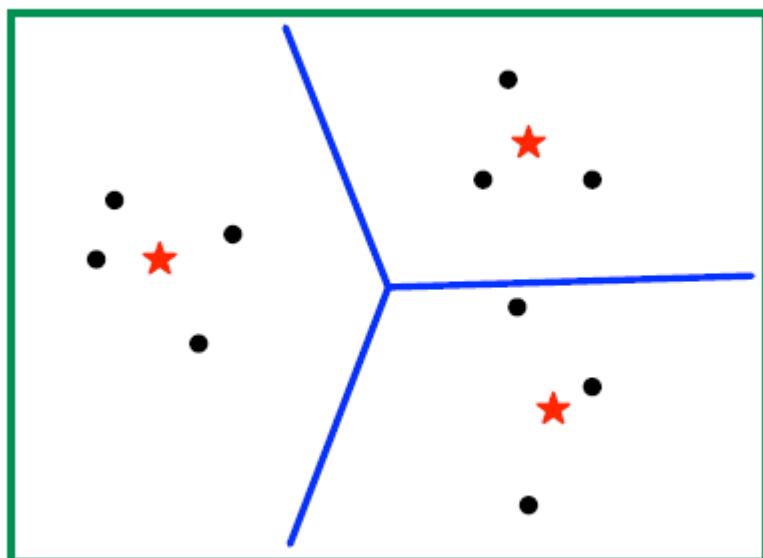
$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



# The $k$ -means optimization problem

- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^p$ ; integer  $k$
- Output: “Centers”, or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:

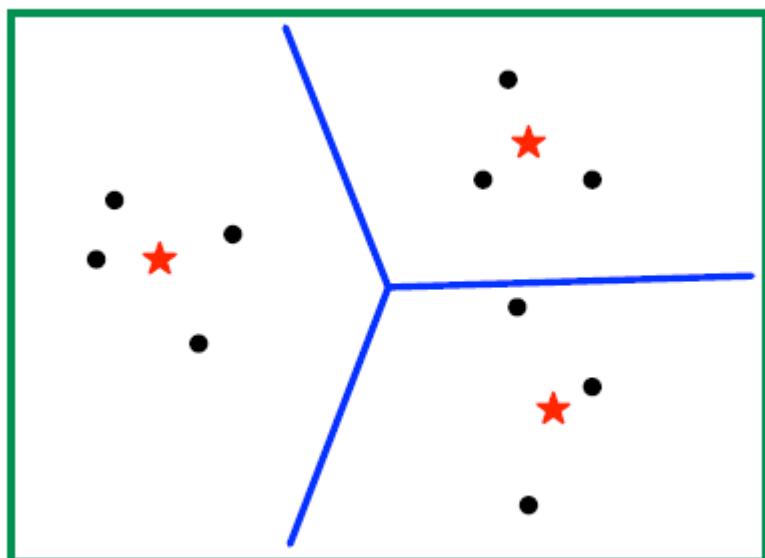
$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$



# The $k$ -means optimization problem

- Input: Points  $x_1, \dots, x_n \in \mathbb{R}^p$ ; integer  $k$
- Output: “Centers”, or representatives,  $\mu_1, \dots, \mu_k \in \mathbb{R}^p$
- Goal: Minimize average squared distance between points and their nearest representatives:

$$\text{cost}(\mu_1, \dots, \mu_k) = \sum_{i=1}^n \min_j \|x_i - \mu_j\|^2$$

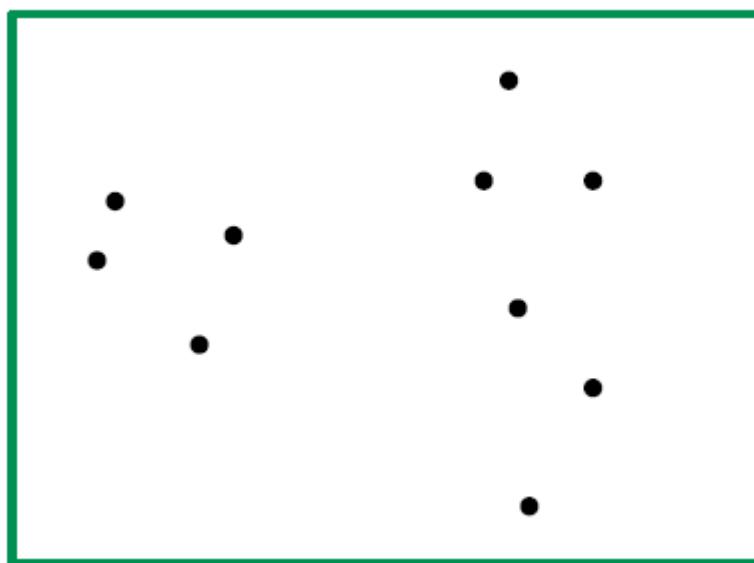


The centers carve  $\mathbb{R}^p$  up into  $k$  convex regions:  $\mu_j$ 's region consists of points for which it is the closest center.

# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

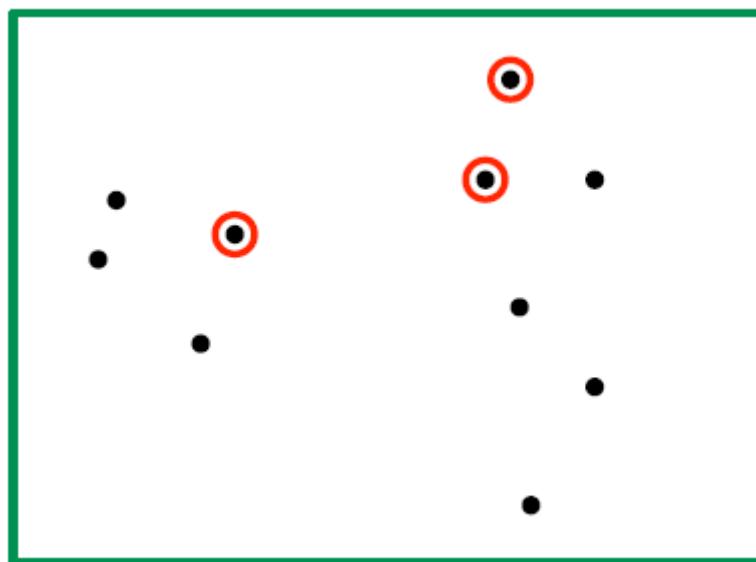
- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

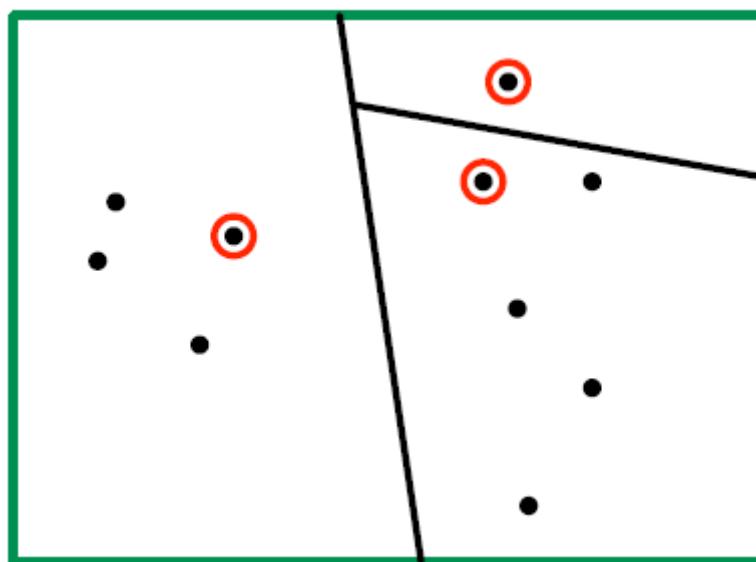
- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

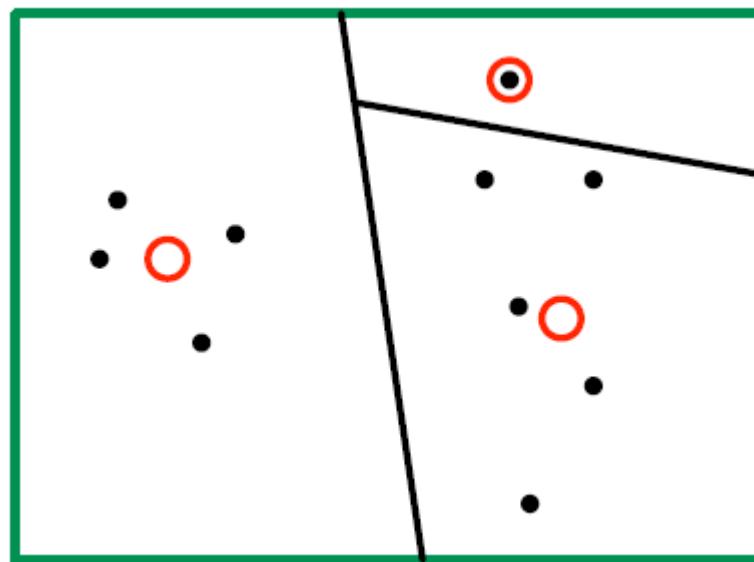
- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

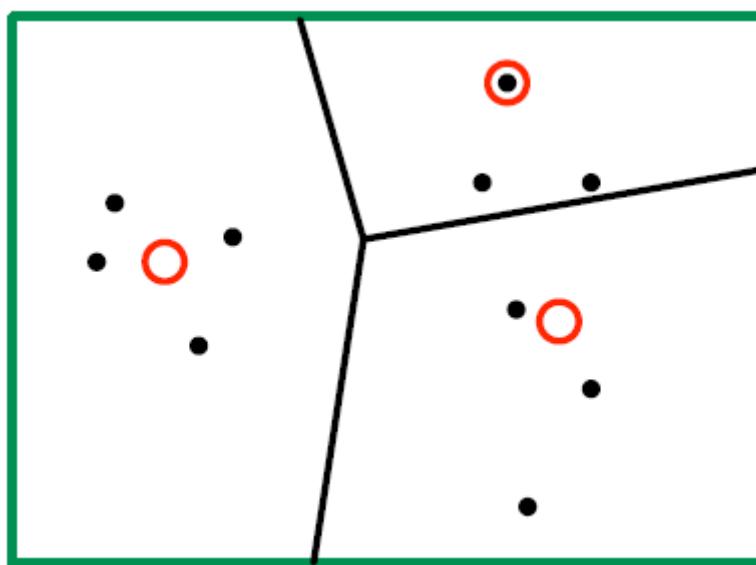
- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

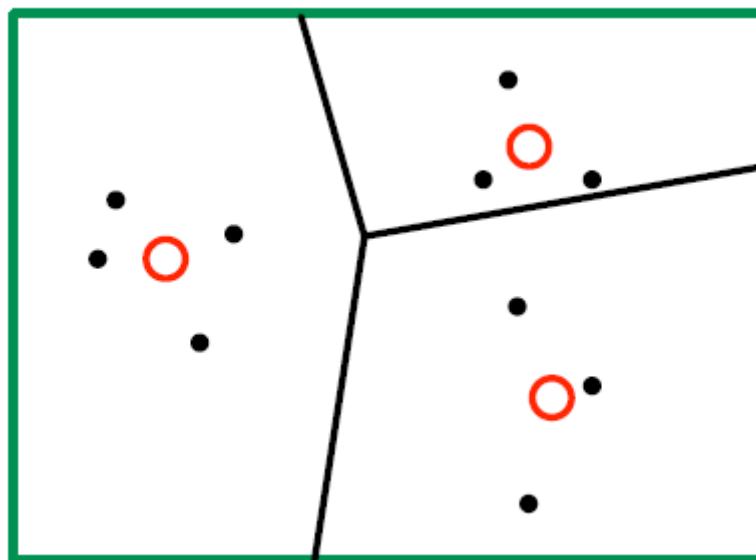
- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

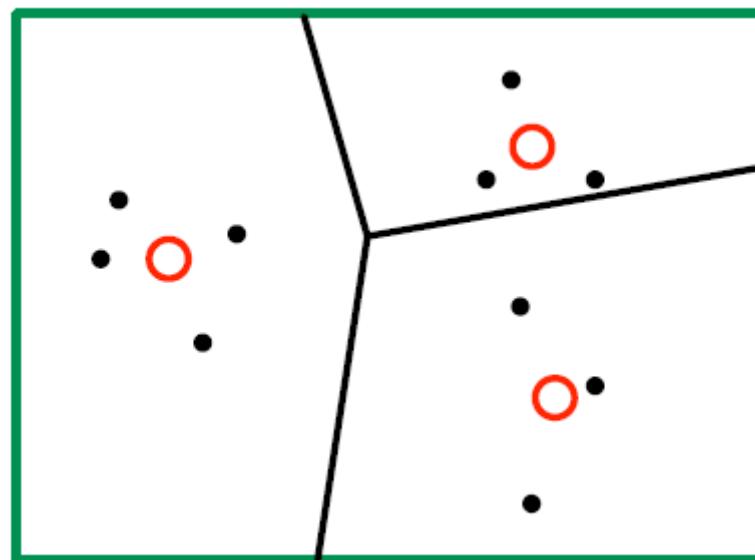
- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



# Lloyd's $k$ -means algorithm

The  $k$ -means problem is NP-hard to solve. The most popular heuristic is called the “ $k$ -means algorithm”.

- Initialize centers  $\mu_1, \dots, \mu_k$  in some manner.
- Repeat until convergence:
  - Assign each point to its closest center.
  - Update each  $\mu_j$  to the mean of the points assigned to it.



Each iteration reduces the cost  $\Rightarrow$  convergence to a local optimum.

# Initializing the $k$ -means algorithm

Typical practice: choose  $k$  data points at random as the initial centers.

# Initializing the $k$ -means algorithm

Typical practice: choose  $k$  data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.

# Initializing the $k$ -means algorithm

Typical practice: choose  $k$  data points at random as the initial centers.

Another common trick: start with extra centers, then prune later.

A particularly good initializer:  **$k$ -means++**

- Pick a data point  $x$  at random as the first center
- Let  $C = \{x\}$  (centers chosen so far)
- Repeat until desired number of centers is attained:
  - Pick a data point  $x$  at random from the following distribution:

$$\Pr(x) \propto \text{dist}(x, C)^2,$$

- where  $\text{dist}(x, C) = \min_{z \in C} \|x - z\|$
- Add  $x$  to  $C$

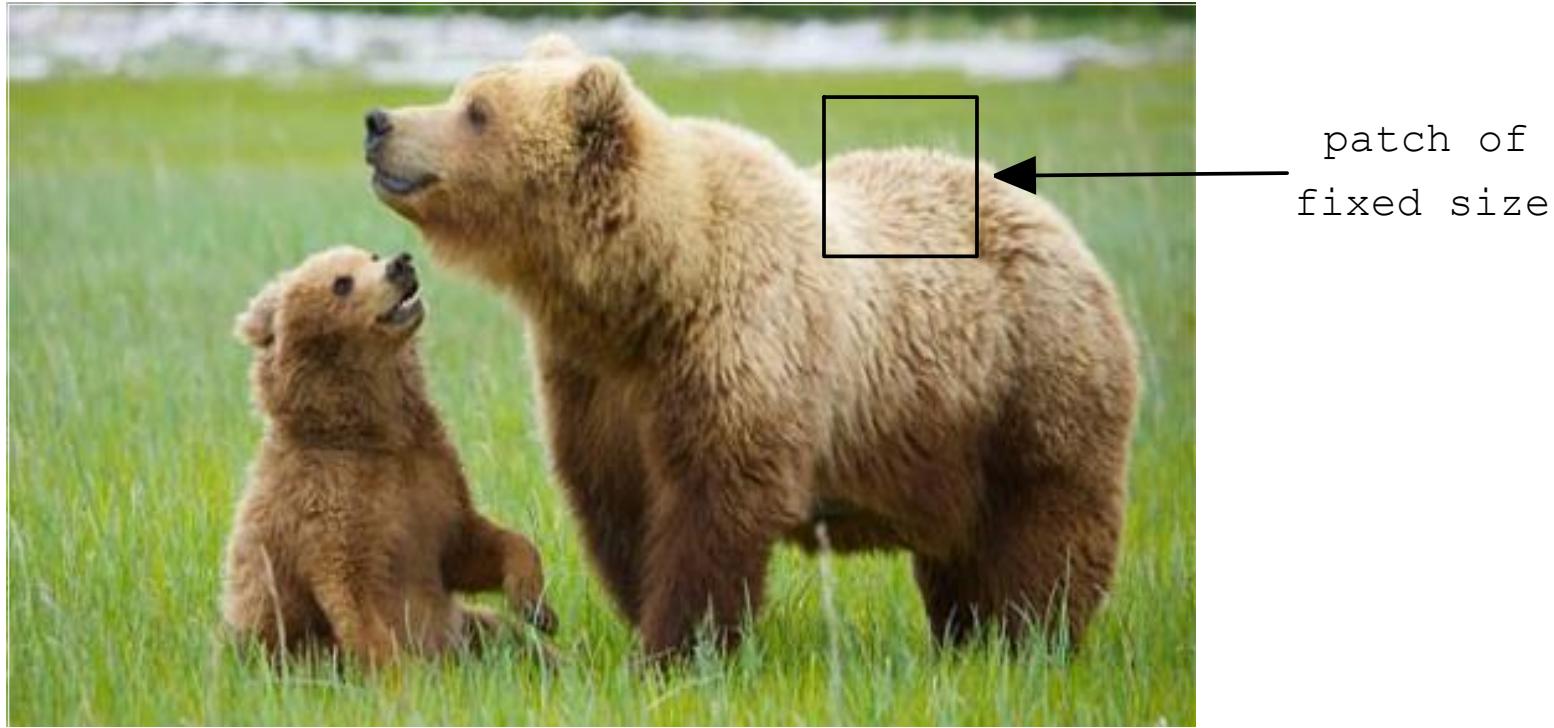
# Representing images using $k$ -means codewords

Given a collection of images, how to represent as fixed-length vectors?



# Representing images using $k$ -means codewords

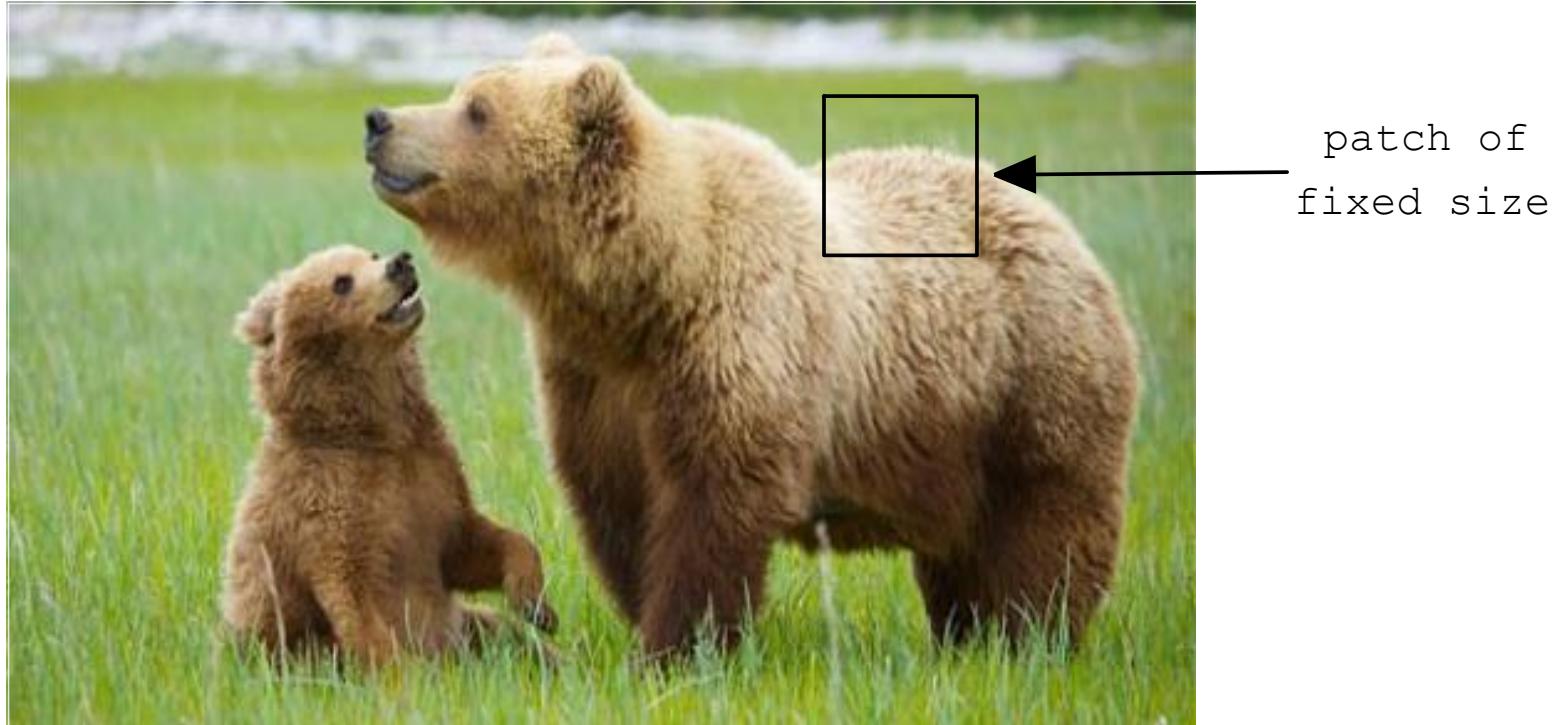
Given a collection of images, how to represent as fixed-length vectors?



- Look at all  $\ell \times \ell$  patches in all images. Extract features for each.
- Run  $k$ -means on this entire collection to get  $k$  centers.
- Now associate any image patch with its nearest center.
- Represent an image by a histogram over  $\{1, 2, \dots, k\}$ .

# Representing images using $k$ -means codewords

Given a collection of images, how to represent as fixed-length vectors?



- Look at all  $\ell \times \ell$  patches in all images. Extract features for each.
- Run  $k$ -means on this entire collection to get  $k$  centers.
- Now associate any image patch with its nearest center.
- Represent an image by a histogram over  $\{1, 2, \dots, k\}$ .

Such data sets are truly enormous.

# Streaming and online computation

**Streaming computation:** for data sets that are too large to fit in memory.

- Make one pass (or maybe a few passes) through the data.
- On each pass:
  - See data points one at a time, in order.
  - Update models/parameters along the way.
- There is only enough space to store a tiny fraction of the data, or a perhaps short summary.

# Streaming and online computation

**Streaming computation:** for data sets that are too large to fit in memory.

- Make one pass (or maybe a few passes) through the data.
- On each pass:
  - See data points one at a time, in order.
  - Update models/parameters along the way.
- There is only enough space to store a tiny fraction of the data, or a perhaps short summary.

**Online computation:** an even more lightweight setup, for data that is continuously being collected.

- Initialize a model.
- Repeat forever:
  - See a new data point.
  - Update model if need be.

# Example: sequential $k$ -means

- ① Set the centers  $\mu_1, \dots, \mu_k$  to the first  $k$  data points
- ② Set their counts to  $n_1 = n_2 = \dots = n_k = 1$
- ③ Repeat, possibly forever:
  - Get next data point  $x$
  - Let  $\mu_j$  be the center closest to  $x$
  - Update  $\mu_j$  and  $n_j$ :

$$\mu_j = \frac{n_j \mu_j + x}{n_j + 1} \quad \text{and} \quad n_j = n_j + 1$$

# $K$ -means: the good and the bad

The good:

- Fast and easy.
- Effective in quantization.

The bad:

- Geared towards data in which the clusters are spherical, and of roughly the same radius.

# $K$ -means: the good and the bad

The good:

- Fast and easy.
- Effective in quantization.

The bad:

- Geared towards data in which the clusters are spherical, and of roughly the same radius.

Is there is a similarly-simple algorithm in which clusters of more general shape are accommodated?

# Mixtures of Gaussians

Idea: model each cluster by a Gaussian:

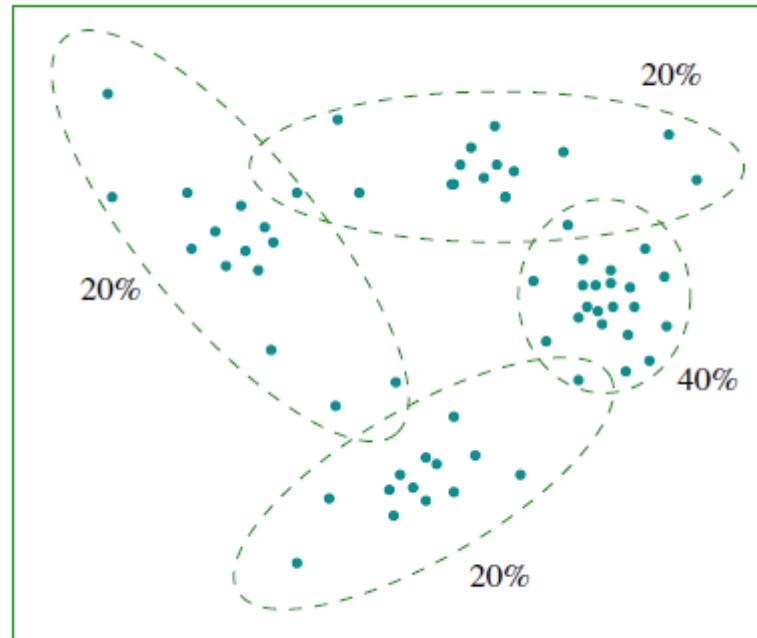
# Mixtures of Gaussians

Idea: model each cluster by a Gaussian:



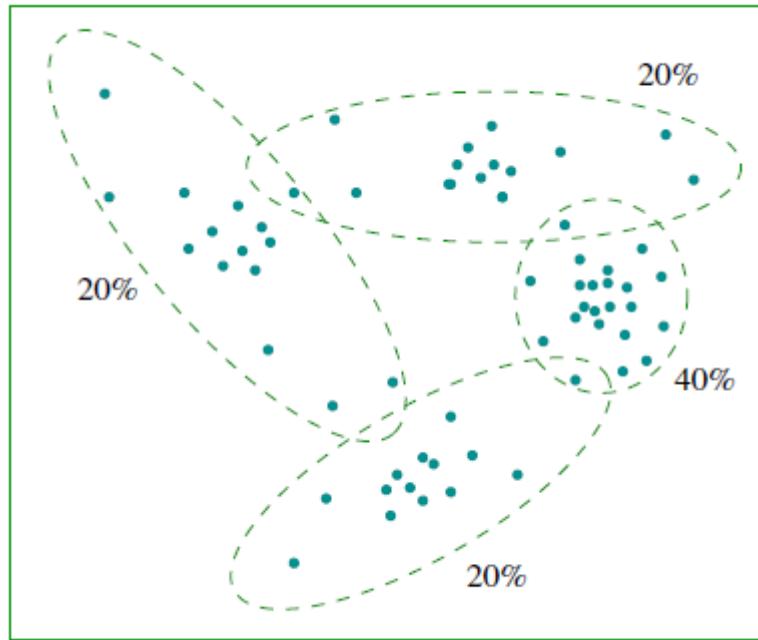
# Mixtures of Gaussians

Idea: model each cluster by a Gaussian:



# Mixtures of Gaussians

Idea: model each cluster by a Gaussian:

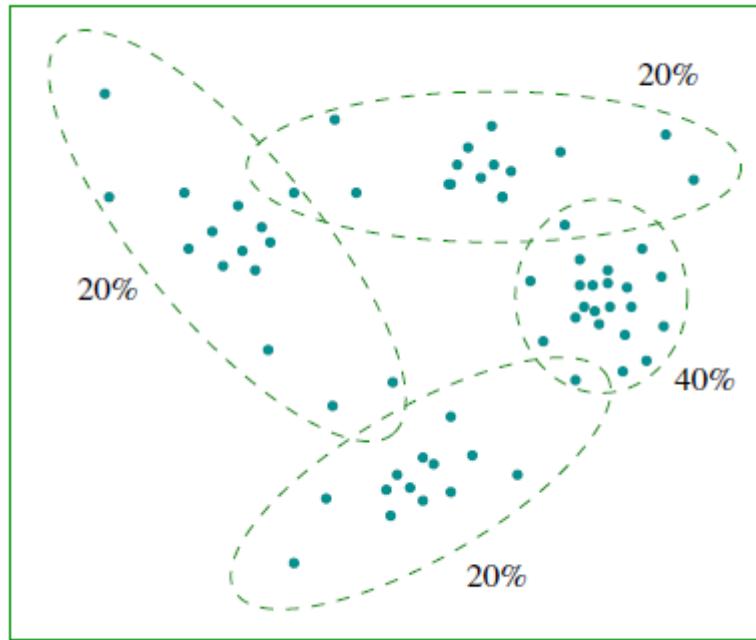


Each of the  $k$  clusters is specified by:

- a Gaussian distribution  $P_j = N(\mu_j, \Sigma_j)$
- a mixing weight  $\pi_j$

# Mixtures of Gaussians

Idea: model each cluster by a Gaussian:



Each of the  $k$  clusters is specified by:

- a Gaussian distribution  $P_j = N(\mu_j, \Sigma_j)$
- a mixing weight  $\pi_j$

Overall distribution over  $\mathbb{R}^p$ : a **mixture of Gaussians**

$$\Pr(x) = \pi_1 P_1(x) + \cdots + \pi_k P_k(x)$$

# The clustering task

Given data  $x_1, \dots, x_n \in \mathbb{R}^P$ , find the maximum-likelihood mixture of Gaussians: that is, find parameters

- $\pi_1, \dots, \pi_k \geq 0$  summing to one
- $\mu_1, \dots, \mu_k \in \mathbb{R}^P$
- $\Sigma_1, \dots, \Sigma_k \in \mathbb{R}^{P \times P}$

to maximize

$$\begin{aligned}\Pr(\text{data} \mid \pi_1 P_1 + \dots + \pi_k P_k) &= \prod_{i=1}^n (P(x_i \mid \pi_1 P_1 + \dots + \pi_k P_k)) \\ &= \prod_{i=1}^n \left( \sum_{j=1}^k P(x_i, z_j \mid \pi_1 P_1 + \dots + \pi_k P_k) \right) = \prod_{i=1}^n \left( \sum_{j=1}^k \pi_j P(x_i \mid z_j) \right) \\ &= \prod_{i=1}^n \left( \sum_{j=1}^k \pi_j P_j(x_i) \right) \\ &= \prod_{i=1}^n \left( \sum_{j=1}^k \frac{\pi_j}{(2\pi)^{P/2} |\Sigma_j|^{1/2}} \exp \left( -\frac{1}{2} (x_i - \mu_j)^T \Sigma_j^{-1} (x_i - \mu_j) \right) \right)\end{aligned}$$

where  $P_j$  is the distribution of the  $j$ th cluster,  $N(\mu_j, \Sigma_j)$ .

# The EM algorithm

- ① Initialize  $\pi_1, \dots, \pi_k$  and  $P_1 = N(\mu_1, \Sigma_1), \dots, P_k = N(\mu_k, \Sigma_k)$  in some manner.
- ② Repeat until convergence:
  - Assign each point  $x_i$  fractionally between the  $k$  clusters:

$$w_{ij} = \Pr(\text{cluster } j \mid x_i) = \frac{\pi_j P_j(x_i)}{\sum_\ell \pi_\ell P_\ell(x_i)}$$

- Now update the mixing weights, means, and covariances:

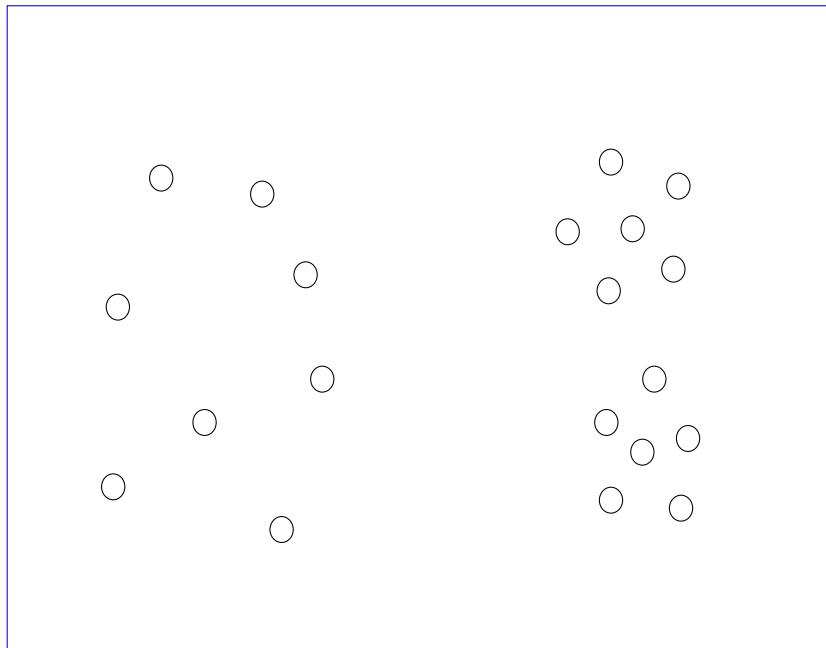
$$\pi_j = \frac{1}{n} \sum_{i=1}^n w_{ij}$$

$$\mu_j = \frac{1}{n\pi_j} \sum_{i=1}^n w_{ij} x_i$$

$$\Sigma_j = \frac{1}{n\pi_j} \sum_{i=1}^n w_{ij} (x_i - \mu_j)(x_i - \mu_j)^T$$

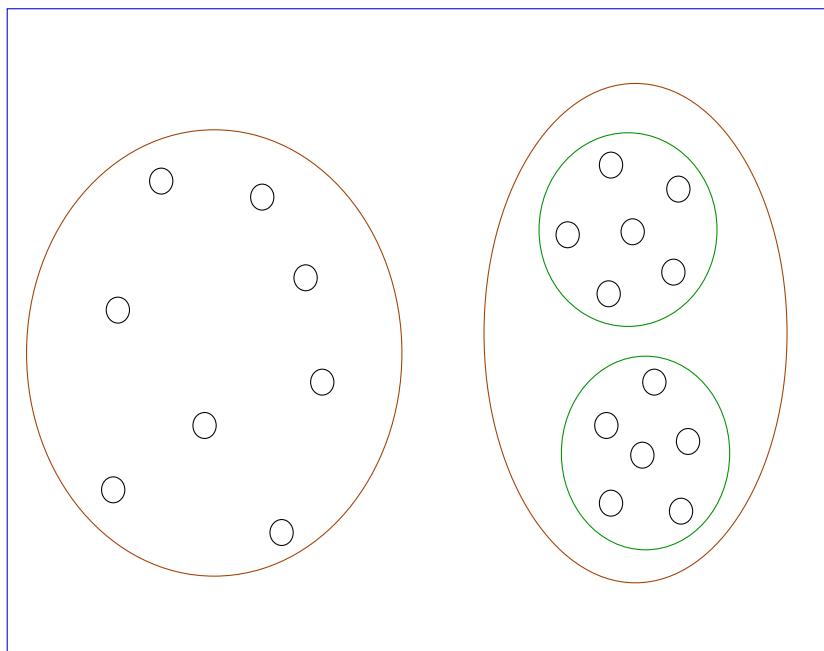
# Hierarchical clustering

Choosing the number of clusters ( $k$ ) is difficult.



# Hierarchical clustering

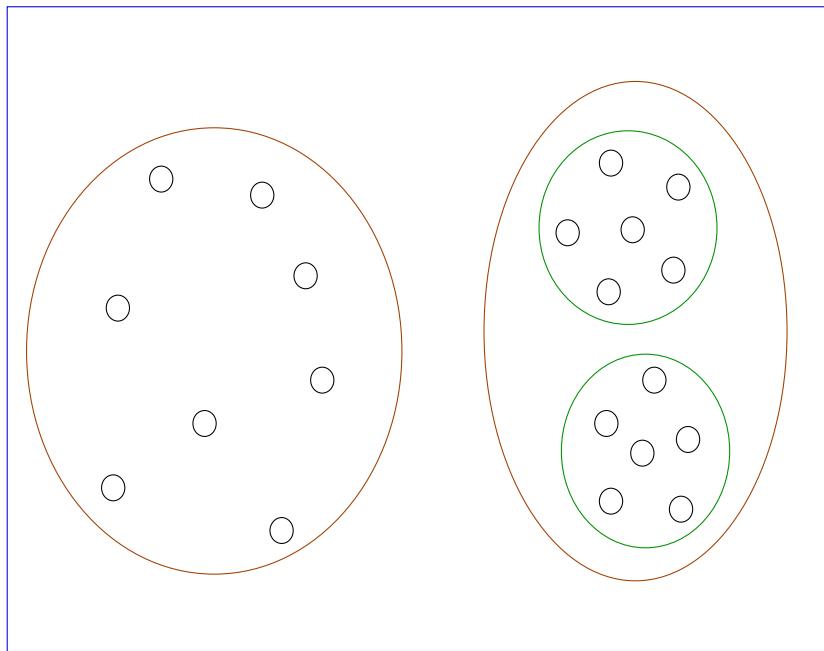
Choosing the number of clusters ( $k$ ) is difficult.



Often there is no single right answer, because of multiscale structure.

# Hierarchical clustering

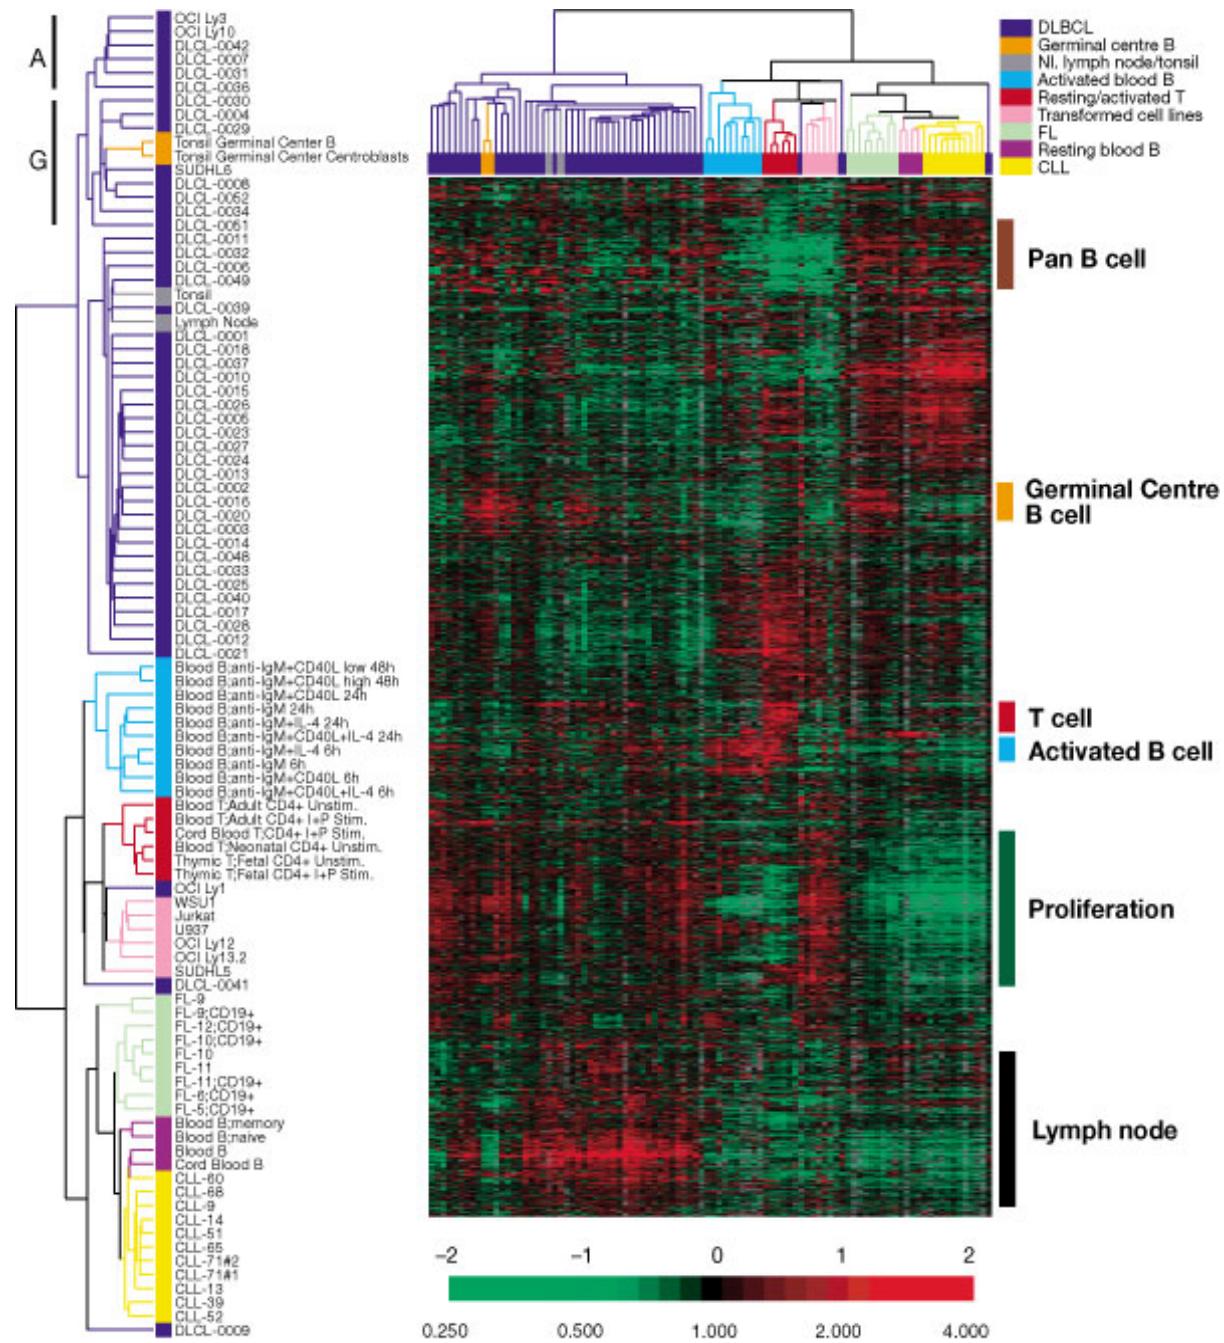
Choosing the number of clusters ( $k$ ) is difficult.



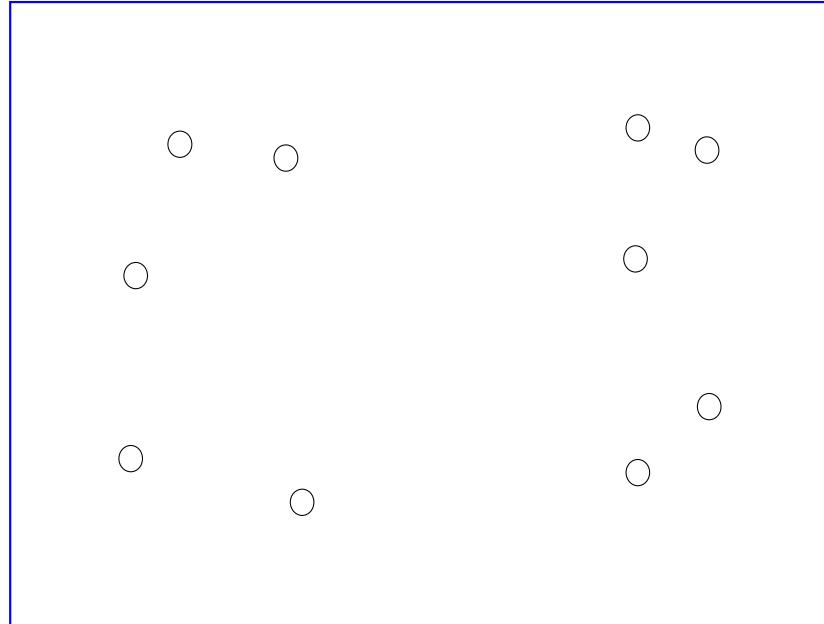
Often there is no single right answer, because of multiscale structure.

Hierarchical clustering avoids these problems.

# Example: gene expression data

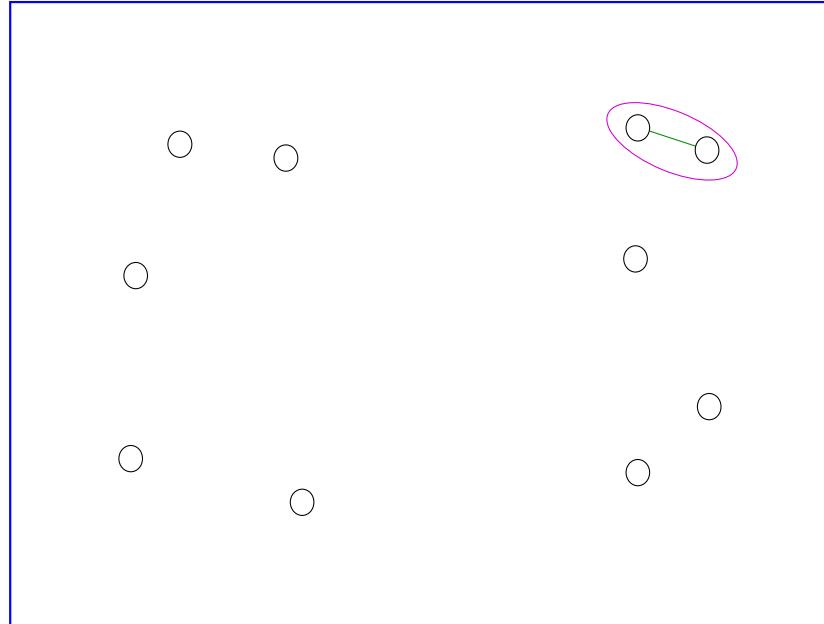


# The single linkage algorithm



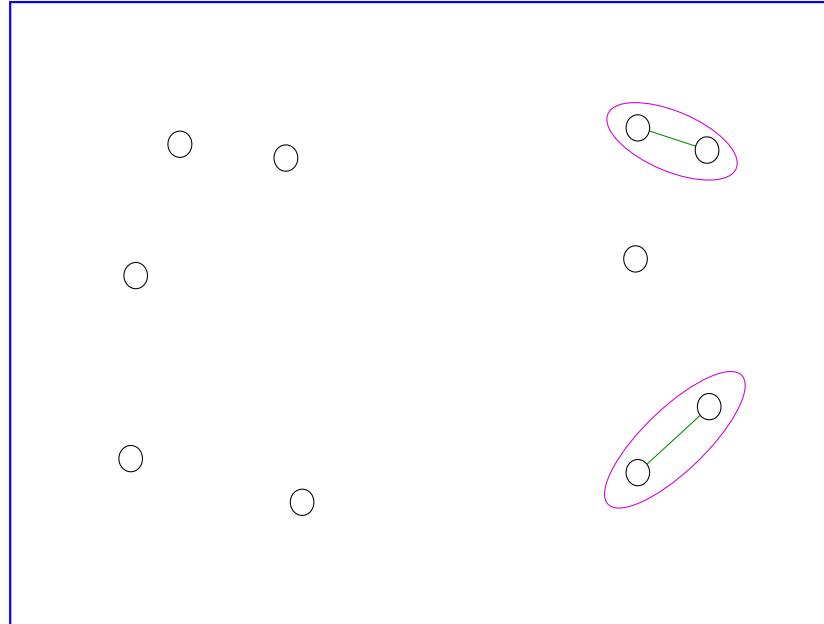
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
  - Disregard singleton clusters

# The single linkage algorithm



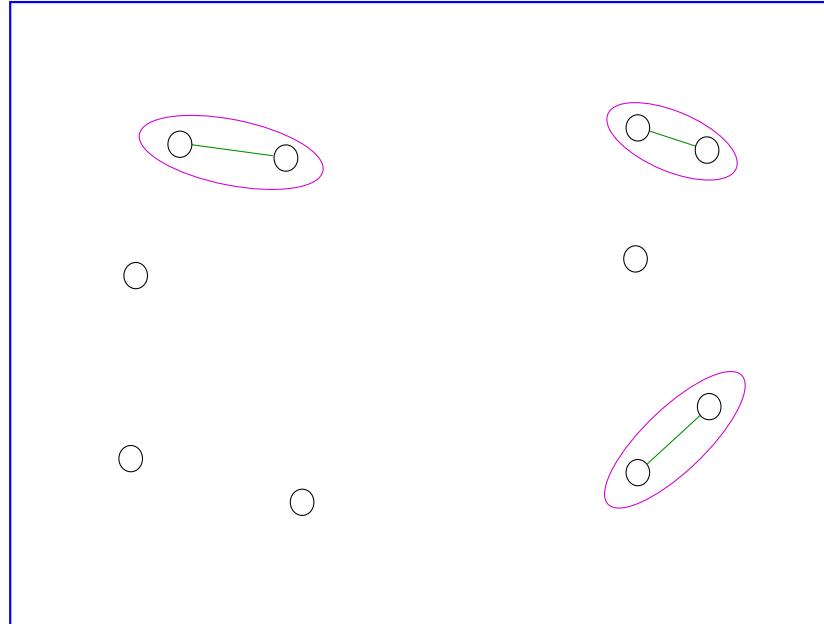
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
  - Disregard singleton clusters

# The single linkage algorithm



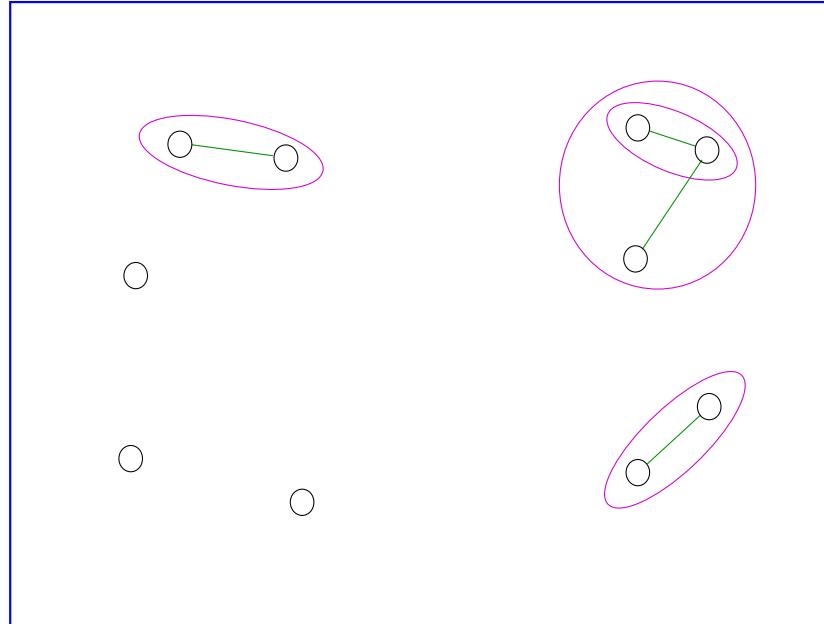
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

# The single linkage algorithm



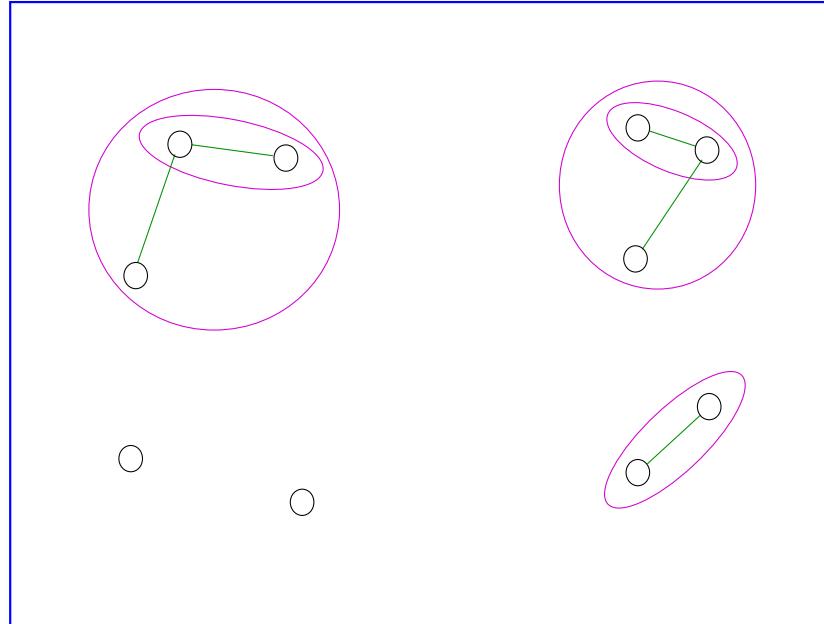
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
  - Disregard singleton clusters

# The single linkage algorithm



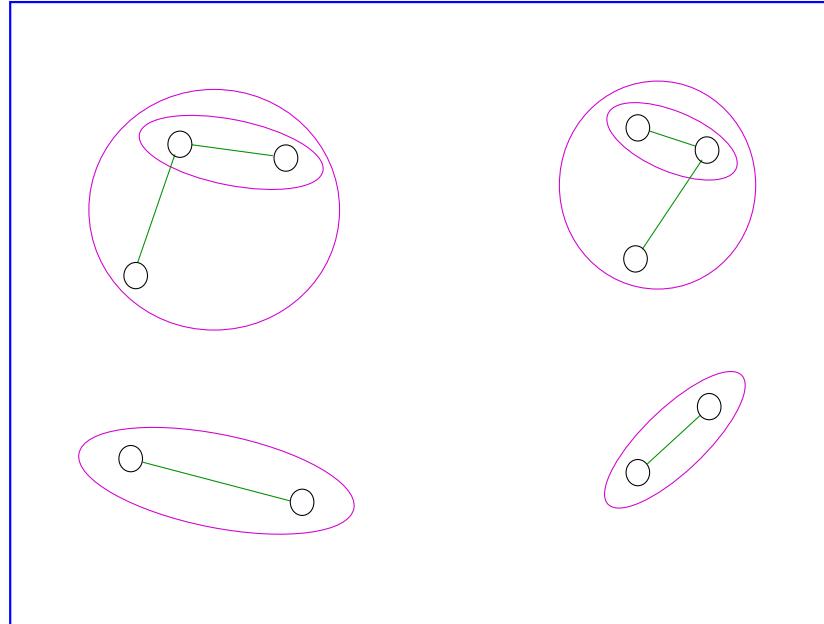
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
  - Disregard singleton clusters

# The single linkage algorithm



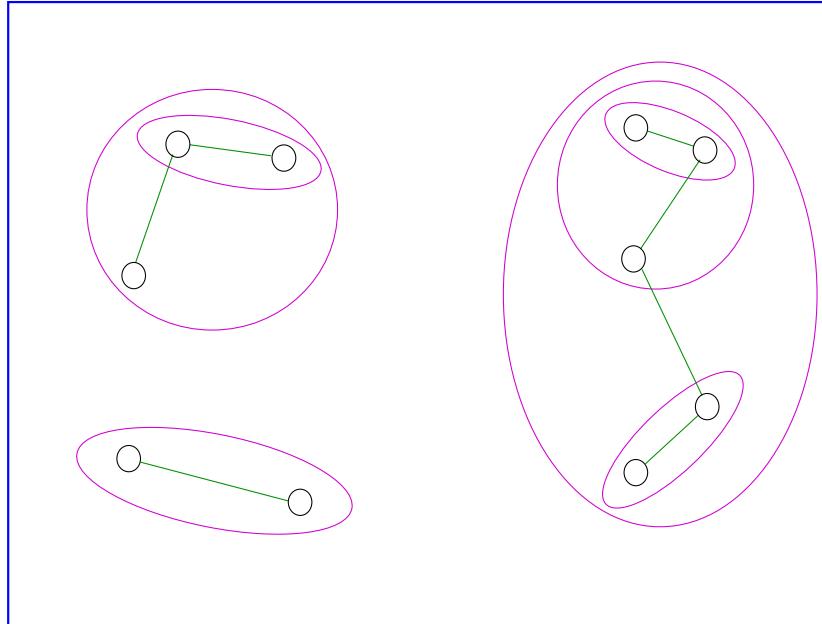
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
  - Disregard singleton clusters

# The single linkage algorithm



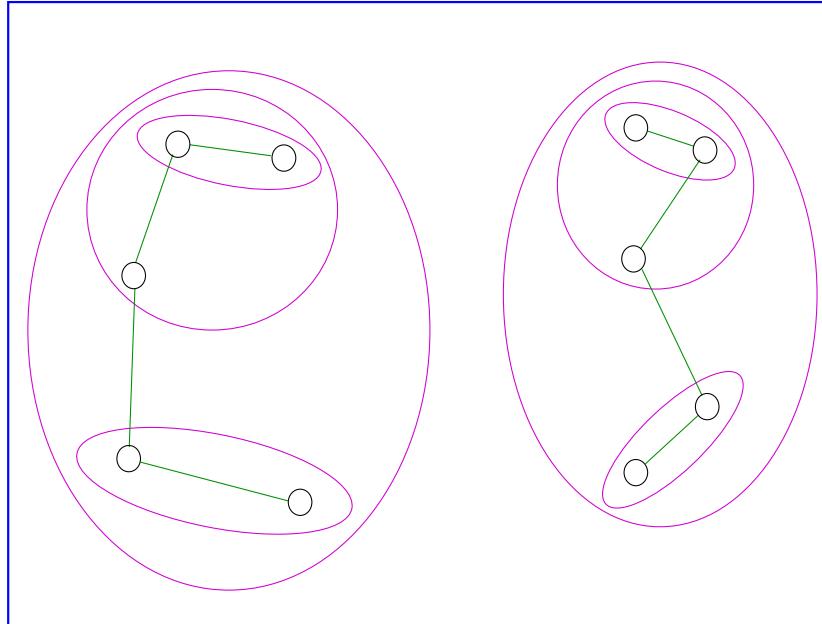
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

# The single linkage algorithm



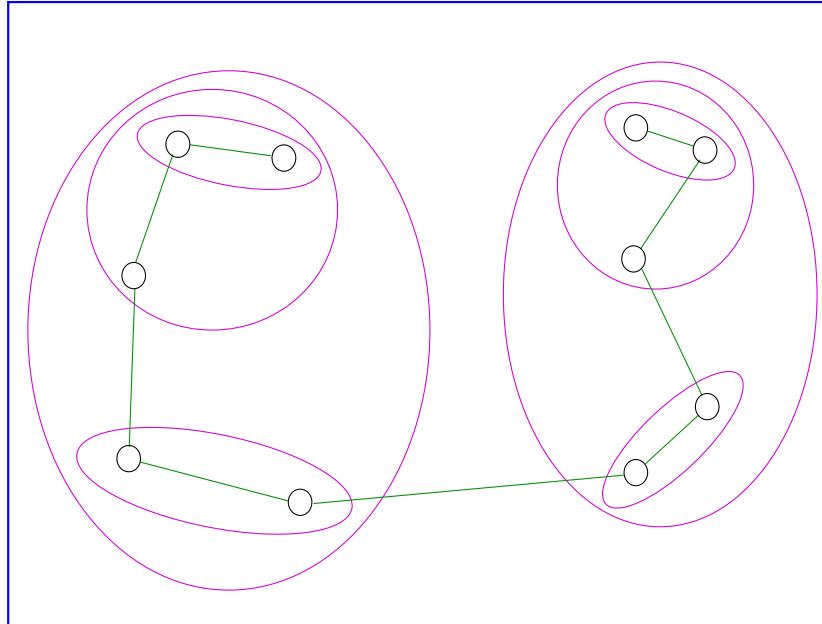
- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

# The single linkage algorithm



- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
  - Disregard singleton clusters

# The single linkage algorithm



- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two clusters with the closest pair of points
- Disregard singleton clusters

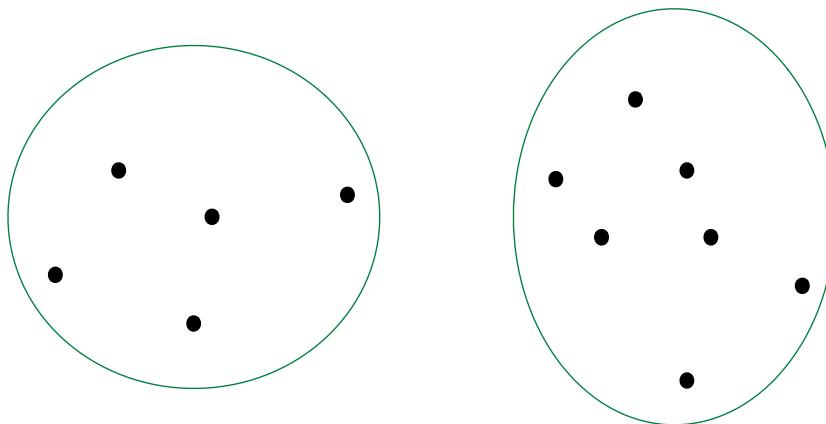
# Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two “closest” clusters

# Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two “closest” clusters

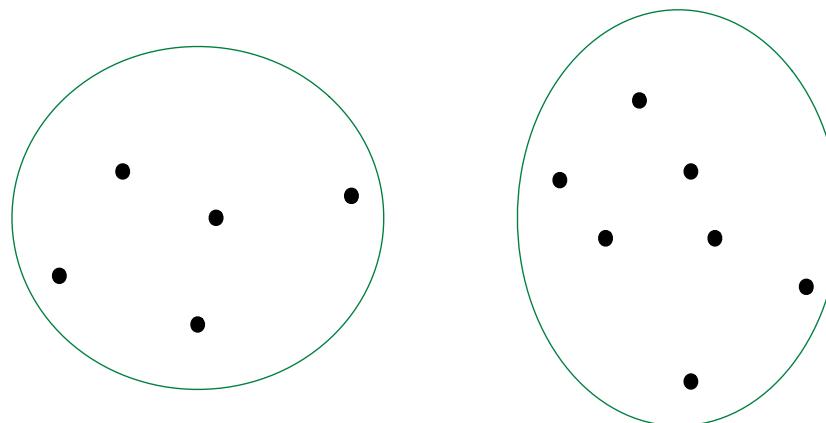
How to measure the distance between two clusters of points,  $C$  and  $C'$ ?



# Linkage methods

- Start with each point in its own, singleton, cluster
- Repeat until there is just one cluster:
  - Merge the two “closest” clusters

How to measure the distance between two clusters of points,  $C$  and  $C'$ ?



- Single linkage

$$\text{dist}(C, C') = \min_{x \in C, x' \in C'} \|x - x'\|$$

- Complete linkage

$$\text{dist}(C, C') = \max_{x \in C, x' \in C'} \|x - x'\|$$

# Average linkage

Three commonly-used variants:

- ① Average pairwise distance between points in the two clusters

$$\text{dist}(C, C') = \frac{1}{|C| \cdot |C'|} \sum_{x \in C} \sum_{x' \in C'} \|x - x'\|$$

- ② Distance between cluster centers

$$\text{dist}(C, C') = \|\text{mean}(C) - \text{mean}(C')\|$$

- ③ Ward's method: the increase in  $k$ -means cost occasioned by merging the two clusters

$$\text{dist}(C, C') = \frac{|C| \cdot |C'|}{|C| + |C'|} \|\text{mean}(C) - \text{mean}(C')\|^2$$

(penalize merging of large clusters)